

Some loops have more than one natural endpoint. Suppose we want to iterate the read-process loop at most ten times. In the example at left below, the "break" statement cannot be criticized, because the loop has two natural endpoints. We could get rid of the "break" by writing the loop as at right below, but the result is longer and harder to read.

<pre>for (int i = 0; i < 10; i++) { s = keybd.readLine(); if (s.length() == 0) { break; } process(s); }</pre>	<pre>int i = 0; do { s = keybd.readLine(); if (s.length() > 0) { process(s); } i++; } while ((i < 10) && (s.length() > 0));</pre>
--	--

There are anti-break zealots who will claim that the loop on the right is the "correct" way to do things. Some of them feel this way because "break" statements are a little bit like the "go to" statements found in some languages like Basic and Fortran (and the "machine language" that microprocessors really execute). "go to" statements allow you to jump to any line of code in the program. It sounds like a good idea at first, but it invariably leads to insanely unmaintainable code. And what happens if you jump to the middle of a loop? Edsger Dijkstra wrote a famous article in 1968 entitled "Go To Statement Considered Harmful", which is part of the reason why many modern languages like Java don't have "go to" statements.

WARNING: It's easy to forget exactly where a "break" statement will jump to. For example, "break" does not jump to the end of the innermost enclosing "if" statement. An AT&T programmer introduced a bug into telephone switching software in a procedure that contained a "switch" statement, which contained an "if" clause, which contained a "break", which was intended for the "if" clause, but instead jumped to the end of the "switch" statement. As a result, on January 15, 1990, AT&T's entire U.S. long distance service collapsed for eleven hours. (That code was actually written in C, but Java's loop syntax and "break" semantics are identical.)

For this reason, Java (unlike C) allows you to attach labels to any kind of enclosing statement, including "if" statements and any group of statements placed { in braces }. Then, "break" can jump to the end of any labeled enclosure that encloses the "break" statement.

```
test:
if (x == 0) {
    loop:
    while (i < 9) {
        stuff: {
            switch(z[i]) {
                case 0: break;           // Jump to statement1
                case 1: break stuff;     // Jump to statement2
                case 2: break loop;      // Jump to statement4
                case 3: break test;      // Jump to statement5
                case 4: continue;        // Jump to location 3
                default: continue loop;  // Jump to location 3
            }
            statement1();
        }
        statement2();
        i++;
        // location 3
    }
    statement4();
}
statement5();
```

The "continue" statement is akin to the "break" statement, except

- (1) it only applies to loops (so you can't write "continue stuff" or "continue test" above), and
- (2) it doesn't necessarily exit the loop; another iteration may commence (if the condition of the "while"/"do"/"for" loop is satisfied).

Unless you want to work for AT&T, I suggest you always use labeled break and continue statements (except perhaps in uncomplicated "switch" statements).

Finally, I told you that "for" loops are identical to certain "while" loops, but there's actually a subtle difference when you use "continue". What's the difference between the following two loops?

<pre>int i = 0; while (i < 10) { if (condition(i)) { continue; } call(i); i++; }</pre>	<pre>for (int i = 0; i < 10; i++) { if (condition(i)) { continue; } call(i); }</pre>
---	---

Answer: when "continue" is called in the "while" loop, "i++" is not executed. In the "for" loop, however, i is incremented at the end of every iteration, even iterations where "continue" is called.

CONSTANTS
=====

Java's "final" keyword is used to declare a value that can never be changed. If you find yourself repeatedly using a numerical value with some "meaning" in your code, you should probably turn it into a "final" constant.

BAD: if (month == 2) {

GOOD: public final static int FEBRUARY = 2; // Usually near top of class.
 ...
 if (month == FEBRUARY) {

Why? Because if you ever need to change the numerical value assigned to February, you'll only have to change one line of code, rather than hundreds.

You can't change the value of FEBRUARY after it is declared and initialized. If you try to assign another value to FEBRUARY, you'll have a compiler error.

"final" is usually used for class variables (static fields), but it can be used for instance variables (non-static fields) and local variables too. It only makes sense for an instance variable to be "final" if the variable is declared with an initializer that calls a method or constructor that doesn't always return the same value.

The custom of rendering constants in all-caps is long-established and was inherited from C.

For any array x, "x.length" is a "final" field.