

CS61B: Lecture 34
Friday, April 19, 2013

Today's reading: Goodrich & Tamassia, Sections 11.3.1 & 11.5.

SELECTION =====

Suppose that we want to find the k th smallest key in a list. In other words, we want to know which item has index j if the list is sorted (where $j = k - 1$). We could simply sort the list, then look up the item at index j . But if we don't actually need to sort the list, is there a faster way? This problem is called `_selection_`.

One example is finding the median of a set of keys. In an array of n items, we are looking for the item whose index is $j = \text{floor}(n / 2)$ in the sorted list.

Quickselect -----

We can modify quicksort to perform selection for us. Observe that when we choose a pivot v and use it to partition the list into three lists I_1 , I_v , and I_2 , we know which of the three lists contains index j , because we know the lengths of I_1 and I_2 . Therefore, we only need to search one of the three lists.

Here's the quickselect algorithm for finding the item at index j - that is, having the $(j + 1)$ th smallest key.

```

Start with an unsorted list  $I$  of  $n$  input items.
Choose a pivot item  $v$  from  $I$ .
Partition  $I$  into three unsorted lists  $I_1$ ,  $I_v$ , and  $I_2$ .
-  $I_1$  contains all items whose keys are smaller than  $v$ 's key.
-  $I_2$  contains all items whose keys are larger than  $v$ 's.
-  $I_v$  contains the pivot  $v$ .
- Items with the same key as  $v$  can go into any of the three lists.
  (In list-based quickselect, they go into  $I_v$ ; in array-based quickselect,
  they go into  $I_1$  and  $I_2$ , just like in array-based quicksort.)
if ( $j < |I_1|$ ) {
    Recursively find the item with index  $j$  in  $I_1$ ; return it.
} else if ( $j < |I_1| + |I_v|$ ) {
    Return the pivot  $v$ .
} else {
    //  $j \geq |I_1| + |I_v|$ .
    Recursively find the item with index  $j - |I_1| - |I_v|$  in  $I_2$ ; return it.
}

```

The advantage of quickselect over quicksort is that we only have to make one recursive call, instead of two. Since we make at most `_one_` recursive call at `_every_` level of the recursion tree, quickselect is much faster than quicksort. I won't analyze quickselect here, but it runs in $\Theta(n)$ average time if we select pivots randomly.

We can easily modify the code for quicksort on arrays, presented in Lecture 32, to do selection. The partitioning step is done exactly according to the Lecture 32 pseudocode for array quicksort. Recall that when the partition stage finishes, the pivot is stored at index "i" (see the variable "i" in the array quicksort pseudocode). In the quickselect pseudocode above, just replace $|I_1|$ with i and $|I_v|$ with 1.

A LOWER BOUND ON COMPARISON-BASED SORTING

=====

Suppose we have a scrambled array of n numbers, with each number from $1 \dots n$ occurring once. How many possible orders can the numbers be in?

The answer is $n!$, where $n! = 1 * 2 * 3 * \dots * (n-2) * (n-1) * n$. Here's why: the first number in the array can be anything from $1 \dots n$, yielding n possibilities. Once the first number is chosen, the second number can be any one of the remaining $n-1$ numbers, so there are $n * (n-1)$ possible choices of the first two numbers. The third number can be any one of the remaining $n-2$ numbers, yielding $n * (n-1) * (n-2)$ possibilities for the first three numbers. Continue this reasoning to its logical conclusion.

Each different order is called a `_permutation_` of the numbers, and there are $n!$ possible permutations. (For Homework 9, you are asked to create a random permutation of maze walls.)

Observe that if $n > 0$,

$$n! = 1 * 2 * \dots * (n-1) * n \leq n * n * n * \dots * n * n * n = n^n$$

and (supposing n is even)

$$n! = 1 * 2 * \dots * (n-1) * n \geq \frac{n}{2} * \frac{n}{2} * \dots * \frac{n}{2} * \frac{n}{2} = \left(\frac{n}{2}\right)^{n/2}$$

so $n!$ is between $(n/2)^{(n/2)}$ and n^n . Let's look at the logarithms of both these numbers: $\log((n/2)^{(n/2)}) = (n/2) \log(n/2)$, which is in $\Theta(n \log n)$, and $\log(n^n) = n \log n$. Hence, $\log(n!)$ is also in $\Theta(n \log n)$.

A `_comparison-based_sort_` is one in which all decisions are based on comparing keys (generally using "if" statements). All actions taken by the sorting algorithm are based on the results of a sequence of true/false questions. All of the sorting algorithms we have studied are comparison-based.

Suppose that two computers run the `_same_` sorting algorithm at the same time on two `_different_` inputs. Suppose that every time one computer executes an "if" statement and finds it true, the other computer executes the same "if" statement and also finds it true; likewise, when one computer executes an "if" and finds it false, so does the other. Then both runs perform exactly the same data movements (e.g. swapping the numbers at indices i and j) in exactly the same order, so they both permute their inputs in `_exactly_` the same way.

A correct sorting algorithm must generate a `_different_` sequence of true/false answers for each different permutation of $1 \dots n$, because it takes a different sequence of data movements to sort each permutation. There are $n!$ different permutations, thus $n!$ different sequences of true/false answers.

If a sorting algorithm asks d true/false questions, it generates $\leq 2^d$ different sequences of true/false answers. If it correctly sorts every permutation of $1 \dots n$, then $n! \leq 2^d$, so $\log_2(n!) \leq d$, and d is in $\Omega(n \log n)$. The algorithm spends $\Omega(d)$ time asking these d questions. Hence,

=====

EVERY comparison-based sorting algorithm takes $\Omega(n \log n)$ worst-case time.

=====

This is an amazing claim, because it doesn't just analyze one algorithm. It says that of the thousands of comparison-based sorting algorithms that haven't even been invented yet, not one of them has any hope of beating $O(n \log n)$ time for all inputs of length n .

Take note that bucket sort is ONLY appropriate when keys are distributed in a small range; i.e. q is in $O(n)$. On Monday we'll study a sorting algorithm called `_radix_sort_` that will fix that problem. The stability of bucket sort will be important for radix sort.